

# Multithreaded Programming and Synchronization

Programming #2

# Task

Write a program that uses multiple threads to speed up the samplesort algorithm

Multithreaded program

You can either use

*Semaphores for synchronization and mutual exclusion, or*

*Mutex locks for mutual exclusion and condition variables for synchronization*

# Objectives

An assessment task related to ILO4 [Practicability]

A learning activity related to ILO 2


To learn how to use POSIX pthreads and semaphore libraries

- have hands-on practice in designing and developing multithreading programs
- create, manage, and coordinate multiple threads in a shared memory environment
- design and implement synchronization schemes for multithreaded processes using semaphores, mutex locks and condition variables.


# Sequential Samplesort

1. Choose  $b$  samples from the data sequence.
2. Sort the  $b$  samples such that  $s_1 < s_2 < s_3 < \dots < s_b$ .
3. Partition the data sequence into  $b + 1$  subsequences (i.e. buckets) based on the samples, such that every element in the  $j^{\text{th}}$  subsequence is greater than  $s_{j-1}$  and smaller than or equal to  $s_j$ . We assume  $s_0 = -\infty$  and  $s_{b+1} = +\infty$ .
4. Sort each subsequence.
5. Concatenate the sorted subsequences to form the sorted data sequence.


42	98	2	31	86	87	5	13	99	44	67	37	17	7	87	3	96	71	40	19	58	13	61	77	11	13	6	81	76	18	24	14	63	59	99	17	36	84	1	48
----	----	---	----	----	----	---	----	----	----	----	----	----	---	----	---	----	----	----	----	----	----	----	----	----	----	---	----	----	----	----	----	----	----	----	----	----	----	---	----




$$x \leq 13$$



$$13 < x \leq 24$$



$$24 < x \leq 58$$



$$58 < x$$

1	2	3	5	6	7	11	13	13	13
---	---	---	---	---	---	----	----	----	----

31	36	37	40	42	44	48	58
----	----	----	----	----	----	----	----

14	17	17	18	19	24
----	----	----	----	----	----

59	61	63	67	71	76	77	81	84	86	87	87	96	98	99	99
----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

How many buckets to use?

How to select samples?

# Parallelizing Samplesort

Samplesort is very well suited and intuitive for parallelization and scaling

The number of buckets equals the number of threads ( $p$ )

Assume there are  $n$  numbers to be sorted and  $n \gg p$

Selection of  $p-1$  samples using oversampling in a distributed manner

# Parallel Samplesort - 5 phases

## Phase 1

Each worker thread gets  $n/p$  numbers

Uses sequential quicksort to sort its 'local' sequence

Each worker selects  $p$  samples from its 'local' sequence at indices

$0, \frac{n}{p^2}, \frac{2n}{p^2}, \dots, \frac{(p-1)n}{p^2}$

} Concurrently

42	98	2	31	86	87	5	13	99	44	67	37	17	7	87	3	96	71	40	19	58	13	61	77	11	13	6	81	76	18	24	14	63	59	99	17	36	84	1	48
----	----	---	----	----	----	---	----	----	----	----	----	----	---	----	---	----	----	----	----	----	----	----	----	----	----	---	----	----	----	----	----	----	----	----	----	----	----	---	----

Thr0										Thr1										Thr2										Thr3									
42	98	2	31	86	87	5	13	99	44	67	37	17	7	87	3	96	71	40	19	58	13	61	77	11	13	6	81	76	18	24	14	63	59	99	17	36	84	1	48

2	5	13	31	42	44	86	87	98	99	3	7	17	19	37	40	67	71	87	96	6	11	13	13	18	58	61	76	77	81	1	14	17	24	36	48	59	63	84	99
---	---	----	----	----	----	----	----	----	----	---	---	----	----	----	----	----	----	----	----	---	----	----	----	----	----	----	----	----	----	---	----	----	----	----	----	----	----	----	----

# Parallel Samplesort - 5 phases

## Phase 2

**Main thread** gathers all  $p \times p$  samples and sorts them

Selects  $p-1$  pivots from the sorted samples at indices:  $p + \lfloor p/2 \rfloor - 1, 2p + \lfloor p/2 \rfloor - 1, \dots, (p-1)p + \lfloor p/2 \rfloor - 1$

2	13	44	87	3	17	40	71	6	13	58	76	1	17	48	63
---	----	----	----	---	----	----	----	---	----	----	----	---	----	----	----

1	2	3	6	13	13	17	17	40	44	48	58	63	71	76	87
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----

1	2	3	6	13	13	17	17	40	44	48	58	63	71	76	87
---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----



# Parallel Samplesort - 5 phases

## Phase 3

Each worker gets the  $p-1$  pivot values

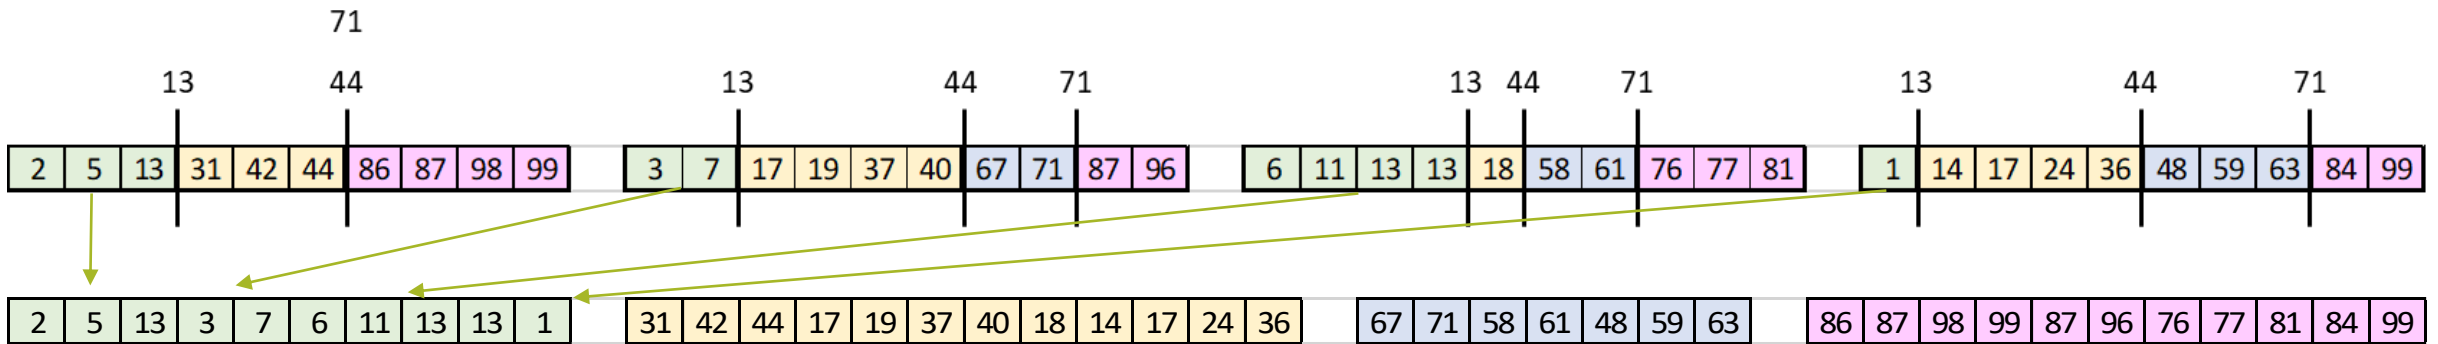
Partitions its 'local' sorted numbers into  $p$  disjoint pieces

Concurrently

## Phase 4

Worker  $i$  keeps its  $i^{th}$  partition and collects from other threads their  $i^{th}$  partitions and merges them into a single sequence

Concurrently



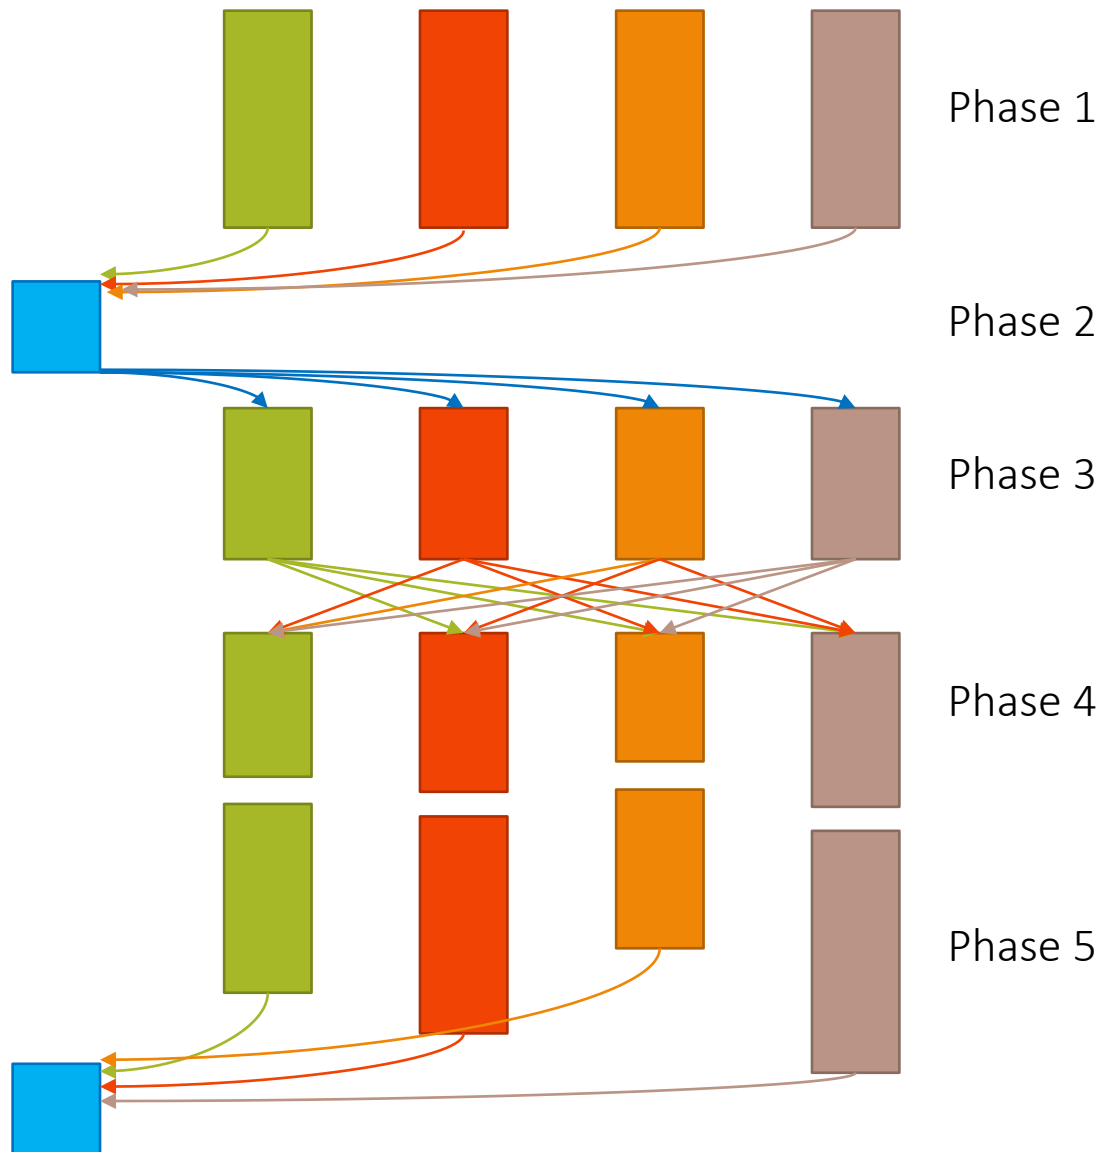
# Parallel Samplesort - 5 phases

## Phase 5

Each worker sorts its sequence } Concurrently  
Main thread gets the final sorted sequence

1	2	3	5	6	7	11	13	13	13		14	17	17	18	19	24	31	36	37	40	42	44		48	58	59	61	63	67	71		76	77	81	84	86	87	87	96	98	99	99
---	---	---	---	---	---	----	----	----	----	--	----	----	----	----	----	----	----	----	----	----	----	----	--	----	----	----	----	----	----	----	--	----	----	----	----	----	----	----	----	----	----	----

1	2	3	5	6	7	11	13	13	13	14	17	17	18	19	24	31	36	37	40	42	44	48	58	59	61	63	67	71	76	77	81	84	86	87	87	96	98	99	99
---	---	---	---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----



Summary

# seqsort.c

```
int main (int argc, char **argv)
{
    long i, j;
    struct timeval start, end;

    if ((argc != 2))
    {
        printf("Usage: seq_sort <number>\n");
        exit(0);
    }

    size = atol(argv[1]);
    intarr = (unsigned int *)malloc(size*sizeof(unsigned int));
    if (intarr == NULL) {perror("malloc"); exit(0); }
```

```
    // set the random seed for generating a fixed random
    // sequence across different runs
    char * env = getenv("RANNUM"); //get the env variable
    if (!env)                       //if not exists
        srand(3230);
    else
        srand(atol(env));

    for (i=0; i<size; i++) {
        intarr[i] = random();
    }
```

```
    // measure the start time
    gettimeofday(&start, NULL);

    // just call the qsort library
    // replace qsort by your parallel sorting algorithm
    // using pthread
    qsort(intarr, size, sizeof(unsigned int), compare);

    // measure the end time
    gettimeofday(&end, NULL);

    if (!checking(intarr, size)) {
        printf("The array is not in sorted order!!\n");
    }

    printf("Total elapsed time: %.4f s\n", (end.tv_sec -
start.tv_sec)*1.0 + (end.tv_usec -
start.tv_usec)/1000000.0);

    free(intarr);
    return 0;
}
```

```
root@1a8d95e9d93a:/home/c3230# gcc seqsort.c -o seqsort
```

```
root@1a8d95e9d93a:/home/c3230# ./seqsort 100000
```

First : 29844

At 25%: 537515580

At 50%: 1074491439

At 75%: 1610064405

Last : 2147482304

Total elapsed time: 0.0223 s

```
root@1a8d95e9d93a:/home/c3230# ./seqsort 100000
```

First : 29844

At 25%: 537515580

At 50%: 1074491439

At 75%: 1610064405

Last : 2147482304

Total elapsed time: 0.0228 s

```
root@1a8d95e9d93a:/home/c3230# export RANNUM='9876'
```

```
root@1a8d95e9d93a:/home/c3230# ./seqsort 100000
```

First : 2351

At 25%: 538105243

At 50%: 1074266763

At 75%: 1611575065

Last : 2147455648

Total elapsed time: 0.0235 s

```
root@1a8d95e9d93a:/home/c3230# export RANNUM='3109'
```

```
root@1a8d95e9d93a:/home/c3230# ./seqsort 100000
```

First : 19740

At 25%: 537365138

At 50%: 1075734945

At 75%: 1607485034

Last : 2147464298

Total elapsed time: 0.0226 s

```
root@1a8d95e9d93a:/home/c3230# unset RANNUM
```

```
root@1a8d95e9d93a:/home/c3230# ./seqsort 100000
```

First : 29844

At 25%: 537515580

At 50%: 1074491439

At 75%: 1610064405

Last : 2147482304

Total elapsed time: 0.0238 s

```
root@1a8d95e9d93a:/home/c3230#
```

# Your task

1. Duplicate a copy of the seqsort.c program and change its name to psort.c
2. Accept one *optional* input argument
  - ./psort <number> [<no\_of\_workers>]
  - default value of no\_of\_workers is 4
3. Implement the parallel Samplesort algorithm by using multiple threads

```
// measure the start time  
gettimeofday(&start, NULL);
```

```
// add the parallel implementation of the samplesort algorithm here
```

```
// measure the end time  
gettimeofday(&end, NULL);
```

4. The random sequence is accessible via the pointer *intarr* before the sorting and we expect you to use the same pointer *intarr* to point to the final sorted sequence

# Measure the multithreaded program

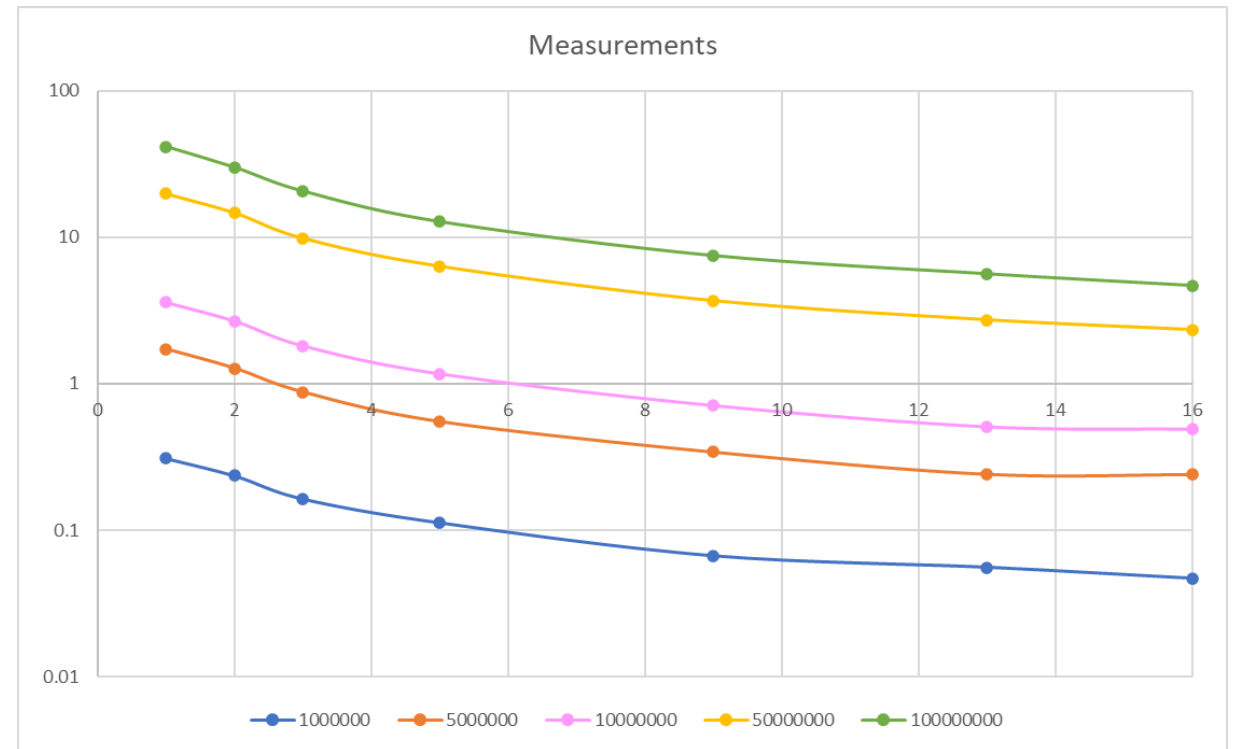
Run the programs on Workbench2 Linux server.

	No. of worker threads	Data size
seqsort	NA	1,000,000, 5,000,000, 10,000,000, 50,000,000, 100,000,000
psort	2, 3, 5, 9, 13, 16	1,000,000, 5,000,000, 10,000,000, 50,000,000, 100,000,000

Complete and submit the measure.xlsx file.

# Sample Measurements

	1000000	5000000	10000000	50000000	100000000
1	0.3087	1.7186	3.5992	19.8209	41.5514
2	0.2349	1.2732	2.6839	14.679	30.0242
3	0.1633	0.8765	1.8161	9.8536	20.6842
5	0.1123	0.5514	1.1674	6.3372	12.7964
9	0.0668	0.3415	0.7075	3.6856	7.5123
13	0.0557	0.241	0.506	2.7264	5.6416
16	0.0469	0.2395	0.4865	2.3305	4.6927





# Computer platform to use

You can develop and test your program on any Linux/Mac/WSL platform.

Finally test and benchmark the programs on Workbench2 server

Your program should be written in C and successfully compiled with gcc

Your submission will be primarily tested under workbench2 server. Make sure that your program can be compiled *without any error*. Otherwise, we have no way to test your submission and thus you may lose *all* the marks

# Submission of Assignment

You should name your program to “psort\_StudentNumber.c” (replace StudentNumber with your HKU student number) and the measurement excel file to measure\_StudentNumber.xlsx

Submit your program and excel file to the course’s Moodle web site (to Programming #2 submission page)

Add your signature in the header of the submitted program and make clear documentation

```
/* **** */
* Filename: psort_3015234567.c
* Student name and Number: Harry Potter 3015234567
* Development platform: WSL2
* Remark: Complete all features
**** */
```

# Grading Criteria

You can use either semaphores or (mutex locks and condition variables) for concurrency and synchronization control. They will be graded without a difference.

Please read the assignment document

<b>Documentation (1 point)</b>	Include necessary documentation to clearly indicate the logic of the program Include required student's info at the beginning of the program
<b>Measurement (1 point)</b>	The performance of the seqsort and psort programs on the Workbench2 server.
<b>psort.c (8 points)</b>	Workers execute in parallel and with any number of worker threads and data sizes
	Correct result
	Threads terminate successfully and allow main thread to collect all sorted sequences
	Main thread wait for all workers to terminate before displaying the results